

ACTL3142: Statistical Machine Learning for Risk and Actuarial Applications

Airbnb Sydney Price Analysis Assignment Part 2

Audrey Chang
Z5627566

1.1 Creating the HAS Variable

The proportion of listings with HAS = 1 is 0.3115. 31.15% of listings are classified as high scoring based on user reviews.

2.1 Data Preparation Cleaning

- Removed irrelevant columns: bathrooms_text, host_since.
- Created host_experience_years = 2025 - host_since.
- Create high_availability: Binary feature based on the median of availability_60.
- Imputed missing values: host_response_time-"Unknown"; host_response_rate and host_acceptance_rate-median values; host_is_superhost-most frequent value (0, not a superhost)

2.2 Feature Engineering

- **Region:** Derived from neighbourhood using domain-specific groupings (East, West, North, South).
- **Property Group:** Categorized property_type into logical groups (e.g., Apartment/Condo, Hotel/Hostel, Guest Accommodation).
- **High Availability:** Binary feature based on whether availability_60 is above the median.
- **Interaction Terms:**
 - host_experience_years × host_is_superhost: Indicates consistent quality over time.
 - price × availability_60: Captures trade-offs between price and availability.
 - room_type × property_group: Captures contextual differences in guest expectations.
- **Log transformation:** Applied to price to handle skewness.
- **Re-leveled categorical variables:** Adjusted reference levels for interpretability.

2.3 Exploratory Analysis and Variable Selection

- Removed avg_score and individual review scores to prevent data leakage.
- Assessed predictive strength of numerical and binary predictors with HAS using point-biserial and phi correlations, respectively. Top features:(See Appendix for correlation)
 - Positively correlated: host_is_superhost, price, longitude, host_experience_years, latitude
 - Negatively correlated: calculated_host_listings_count, instant_bookable, availability_60
- Visualized categorical predictor relationships with HAS using stacked bar plots. (See appendix)

2.4 Train-Test Split

- Random 75%/25% split for training and testing.

2.5 Model Comparison and Justification(See Appendix for Model Comparison and ROC curves and AUC scores)

2.6 Model Selection and Justification(See Appendix for model summary output.)

The logistic regression model was selected as the final model for the following reasons:

Criterion	Justification
Balanced Predictive Power	Achieved the highest AUC among all models with balanced accuracy, recall & F1-score.
High Interpretability	Provides transparent coefficients, ideal for stakeholder interpretation.
Strategic Feature Inclusion	Includes key predictors like host_is_superhost, log_price, and meaningful interactions.
Justifiable Feature Selection	Selected using stepwise selection of lowest AIC(6421.5) and EDA insights, ensuring model parsimony and relevance.
Best Recall Performance	Recall of 60.1%, the highest among all models, ensures accurate detection of high-quality listings.
Most Competitive Overall	Although Ridge had slightly higher accuracy, its recall was significantly lower.

3.1 Model Evaluation

- **Model Fit:** The model's Null Deviance (7145.5) and Residual Deviance (6371.5) show a substantial reduction, indicating a good fit. The AIC (6421.5) suggests the model is relatively efficient. Most predictors are statistically significant, reinforcing the model's strength.

- **Model Coefficients:** Most of the model's predictors are statistically significant and their coefficients are directionally sensible, meaning they align with expectations based on domain knowledge.

3.1.1 Diagnostics & Validation

- **Confusion Matrix:** Shows balanced classification across both HAS classes, with no major imbalance.
- **ROC Curve & AUC:** AUC = 0.722 indicates solid discrimination power between HAS = 1 and HAS = 0.
- **Model Coefficients:** Most predictors are statistically significant and directionally sensible (details in Q4).

3.1.2 Residual & Deviance Diagnostics

- Deviance residuals show slight skew above zero and mild heteroscedasticity, especially at low predicted probabilities.
- A few residuals exceed ± 2 , suggesting moderate outliers.
- No strong curvature or funnel shape in plots, supporting adequate model fit.

3.1.3 Influence Diagnostics:

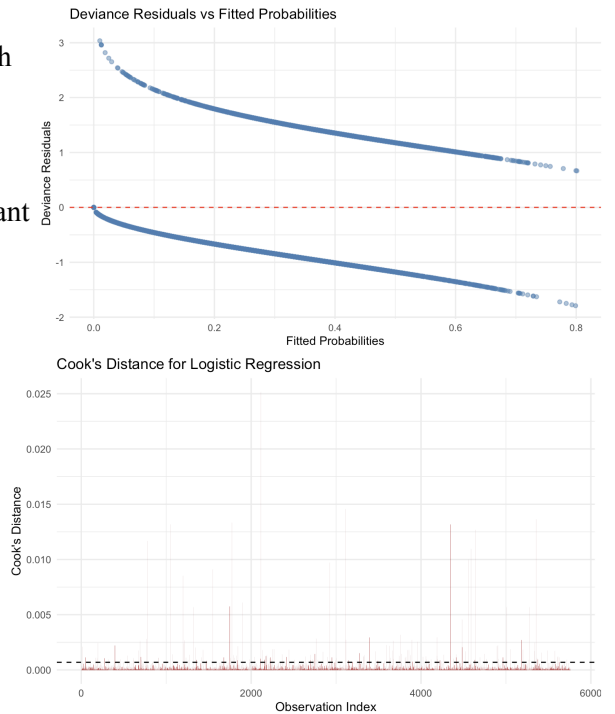
- Cook's Distance values are mostly below 0.025, with no major leverage issues.
- Minor spikes (around index 2000 and beyond) are non-distorting.

3.1.4 Potential Limitations

Issue	Explanation
Linearity Assumption	Logistic regression assumes log-odds are linearly related to predictors. This may oversimplify effects (e.g., price or geography).
Interaction Overfitting	Sparse categories in some room-type \times property-group interactions could lead to overfitting.
Recall at 60.1%	Although balanced, the model still misses some HAS=1 listings
Limited Flexibility	Logistic regression may underperform compared to nonlinear models like Random Forest or XGBoost in capturing complex patterns.
Business Impact of False Negatives	With a recall of 60.1%, ~40% of truly high-rated listings are missed. For Airbnb hosts and investors, these false negatives may lead to lost visibility, underpricing, and fewer bookings — making recall improvement a key business priority.

4.1 Insights for a Sydney Investor with HAS=0 Listings

Recommendation	Model Evidence	Strategic Action
Become a Superhost	host_is_superhost, +0.5361, $p < 2e-16$	Focus on improving host responsiveness, enhancing reviews, and avoiding cancellations to elevate the profile.
Align Price with Quality	log_price, +0.7267, $p < 2e-16$	Invest in premium property features and avoid underpricing when offering high-quality value.
Avoid Overextending	calculated_host_listings_count, -0.0181, $p < 2e-16$	Limit the number of listings to maintain high-quality, personalized service for guests.
Limit Excessive Availability	high_availability, -0.3465, $p < 0.001$	Consider offering exclusive or time-limited booking windows to create a sense of exclusivity.



Be Cautious with Instant Book	instant_bookable, -0.4813, $p < 0.001$	Allow for guest communication before booking whenever possible to ensure both parties' expectations are clear.
Favor South Sydney	regionSouth, +0.3523, $p = 0.0267$	Target regions in South Sydney, where listings show a positive regional effect.
Consider Eastern Suburbs	longitude, +0.8272, $p = 0.049$	Focus on properties near beaches or the CBD (e.g., Bondi), as these areas tend to generate higher ratings.

Room Type × Property Group Insights

-Top Performing Combinations

Private rooms generally perform well when paired with home-like or distinctive property types:

- Private room × Guest Accommodation (+1.192, $p < 0.001$), House/Townhouse/Villa (+1.008, $p < 0.001$), and Unique Stays (+1.279, $p < 2e-16$) score highly—guests value privacy, homeliness, and unique experiences.
- Private room × Apartment/Condo (+1.027, $p = 0.0029$) also performs strongly, likely due to affordability and convenience.

Entire home options also see success when the property offers charm or hospitality focus:

- Entire home × Guest Accommodation (+0.8555, $p < 0.001$) and Cottage/Cabin/Bungalow (+0.7475, $p = 0.002$) attract guests seeking either professional hosting or rustic, private retreats.

-Underperforming Combination

- Private room × Hotel/Hostel (−1.887, $p = 0.0015$) performs poorly—likely because guests booking private rooms expect more personalization than standard hotel settings provide.

5.1 Data Preprocessing & Feature Engineering, used the cleaned dataset from Q2. Key steps included:

- Log Transformation on response: Applied to price to reduce skewness; used `log_price` in all models.
- Review Score Proxy: Kept `avg_score` for review-related insights.
- Engineered Features
 - **Activity Duration:** `days_between_reviews` = last - first review.
 - **PCA:** Applied to accommodates, beds, bedrooms, and bathrooms to reduce dimensionality and multicollinearity; retained 1st PC (`PC_capacity`). (Based on correlation plot)

5.2 Feature Selection (See Appendix for Random Forest's %IncMSE plot)

- Used Random Forest's %IncMSE metric to assess variable importance and identify predictive strength.
- Performed cutoff sensitivity analysis using top 10, 15, and 20 features to evaluate trade-offs.
- Selected top 20 predictors to reduce dimensionality and multicollinearity → `df_reduced` used in all models

5.3 Model Development(See Appendix for model summary)

Split: 60% train, 20% validation, 20% test.

5.3.1 Random Forest

Captured non-linear effects and variable interactions. Hyperparameter grid search included:

- `mtry`: 10, 15, 20 and `ntree`: 100, 300, 500

Trained models on training data; best combination selected using validation RMSE and MSE.

5.3.2 Lasso Regression

- Used `alpha = 1`. Design matrix encoded categorical vars.
- 10-fold CV on train set to choose `lambda.min` (0.000197037).
- Final model trained on full train set. Provided sparse feature selection.

5.3.3 Generalized Linear Model

- Gaussian with identity link on `log_price`.
- Transparent baseline. Residual Deviance = 595.22, AIC = 3720.2.

5.3.4 XGBoost

- Strong tabular data performance. Used:
 - `eta = 0.1`, `max_depth = 6`, `objective = reg:squarederror`

- Early stopping CV to select best rounds. Final model trained accordingly.
- Feature Importance: Key contributors include PC_capacity, longitude, room/property interactions, latitude, etc

5.3.5 Gradient Boosted Machines

- The log-transformed target was back-transformed prior to modeling.
- A model was trained using 1,000 trees, a learning rate (shrinkage) = 0.01, and depth = 4.
- 5-fold cross-validation was employed internally to determine the optimal number of boosting iterations.

5.4 Model Evaluation Reported RMSE on original price scale ($\exp(\log_price)$) (See appendix)

5.5 Final Model Justification, Final Choice: XGBoost


- Predictive performance: XGBoost had the lowest test MSE, our primary criterion.
- Flexibility: It captures non-linear relationships and high-order interactions automatically.
- Feature importance: Though interpretation is not required, it provides insight into what vars drive predictions.
- Robustness: Through regularization and early stopping, overfitting was mitigated.

6.1 Strengths of the Final Model

- High Predictive Accuracy: XGBoost achieved the lowest RMSE on both validation and test sets, excelling in capturing non-linear data relationships.
- No Overfitting: Consistent validation (\$107.8) and test (\$107.71) RMSEs indicate low overfitting risk.
- Robust Feature Engineering: Applied log transformation, PCA on capacity variables, and domain-specific features.

6.2 Model Diagnostics & Validation:

- Residual Analysis: Standardized residuals scatter around zero, indicating no major model misspecification
- 5-Fold Cross-Validation with Early Stopping: Prevented overfitting.
- Mild Heteroscedasticity: Observed for high-priced listings (> \$500), expected in pricing contexts.



6.3 Model Limitations

- Black-Box Nature: XGBoost lacks transparency, limiting its interpretability for stakeholders needing explainable models.
- Heteroscedasticity: Increased residual variance at higher price points reduces precision for luxury listings.
- Lack of Temporal Features: Seasonal and event-based pricing patterns were not captured, likely affecting real-world Airbnb pricing.
- No Segmentation: The model assumes a global pricing strategy, missing potential segmentation (e.g., business vs. leisure, luxury vs. budget).

7.1 Improve the Best Model

- **Expand Cross-Validation and Hyperparameter Tuning:** Use 10-fold cross-validation and a wider grid to improve stability and unlock better parameter combinations. Bayesian optimization could enhance tuning efficiency.
- **Implement Ensemble Stacking:** Combine XGBoost with Random Forest or GBM to leverage each model's strengths, reducing bias and variance for better performance.
- **Enhance Feature Selection and Dimensionality Reduction:** Extend PCA and L1 regularization to reduce overfitting by eliminating weak or collinear predictors.
- **Engineer Nonlinear Transformations:** Add splines or polynomial features for continuous predictors (e.g., latitude, host_experience_years) to improve signal representation and interpretability.
- **Deepen Residual Analysis and Subgroup Diagnostics:** Explore residuals by subgroups (e.g., region, room type) and error distribution plots to identify misfit segments and guide further feature engineering.
- **Incorporate Clustering for Latent Segments:** Use clustering (e.g., K-means, hierarchical) to define submarkets, allowing the model to adapt to distinct market segments for improved predictions.

8.1 Final XGBoost model and preprocessing pipeline were applied to the Q8 test set to generate predictions.

Appendix

Model Summaries

2.5

Logistic Regression Stepwise Selection

Start: AIC=6422.71

HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
high_availability + longitude + latitude + instant_bookable +
region + price_availability_interaction + room_type_property_group_interaction

	Df	Deviance	AIC
- latitude	1	6371.5	6421.5
- price_availability_interaction	1	6372.7	6422.7
<none>		6370.7	6422.7
- longitude	1	6374.8	6424.8
- region	3	6386.1	6432.1
- high_availability	1	6388.6	6438.6
- instant_bookable	1	6415.4	6465.4
- log_price	1	6439.7	6489.7
- host_is_superhost	1	6441.3	6491.3
- room_type_property_group_interaction	14	6546.5	6570.5
- calculated_host_listings_count	1	6550.9	6600.9

Step: AIC=6421.54

HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
high_availability + longitude + instant_bookable + region +
price_availability_interaction + room_type_property_group_interaction

	Df	Deviance	AIC
- price_availability_interaction	1	6373.4	6421.4
<none>		6371.5	6421.5
+ latitude	1	6370.7	6422.7
- longitude	1	6375.5	6423.5
- region	3	6388.1	6432.1
- high_availability	1	6389.2	6437.2
- instant_bookable	1	6416.2	6464.2
- log_price	1	6440.6	6488.6
- host_is_superhost	1	6442.9	6490.9
- room_type_property_group_interaction	14	6549.8	6571.8
- calculated_host_listings_count	1	6552.4	6600.4

Step: AIC=6421.39

HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
high_availability + longitude + instant_bookable + region +
room_type_property_group_interaction

	Df	Deviance	AIC
<none>		6373.4	6421.4
+ price_availability_interaction	1	6371.5	6421.5
+ latitude	1	6372.7	6422.7
- longitude	1	6377.3	6423.3
- region	3	6389.2	6431.2
- high_availability	1	6417.3	6463.3
- instant_bookable	1	6417.8	6463.8
- host_is_superhost	1	6445.1	6491.1
- log_price	1	6463.8	6509.8
- room_type_property_group_interaction	14	6550.9	6570.9

```

- host_is_superhost      1      6463.8  6509.8
- log_price              1      6463.8  6509.8
- room_type_property_group_interaction 14  6550.0  6570.0
- calculated_host_listings_count      1  6553.3  6599.3

```

```

Call:
glm(formula = HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
     high_availability + longitude + instant_bookable + region +
     room_type_property_group_interaction, family = binomial,
     data = train)

```

```

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.7773   -0.8863   -0.6168    1.1264    3.0267

```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.299e+02	6.336e+01	-2.051	0.040290 *
host_is_superhost	5.374e-01	6.336e-02	8.483	< 2e-16 ***
log_price	6.527e-01	6.935e-02	9.412	< 2e-16 ***
calculated_host_listings_count	-1.800e-02	1.791e-03	-10.045	< 2e-16 ***
high_availability1	-4.181e-01	6.341e-02	-6.594	4.28e-11 ***
longitude	8.304e-01	4.202e-01	1.976	0.048121 *
instant_bookable	-4.796e-01	7.318e-02	-6.554	5.61e-11 ***
regionEast	-1.643e-02	1.431e-01	-0.115	0.908563
regionNorth	2.084e-01	1.483e-01	1.405	0.159976
regionSouth	3.477e-01	1.580e-01	2.201	0.027753 *
room_type_property_group_interactionHotel room.Apartment/Condo	2.019e-01	8.794e-01	0.230	0.818441
room_type_property_group_interactionPrivate room.Apartment/Condo	1.008e+00	3.440e-01	2.930	0.003392 **
room_type_property_group_interactionEntire home/apt.Cottage/Cabin/Bungalow	7.408e-01	2.420e-01	3.061	0.002203 **
room_type_property_group_interactionPrivate room.Cottage/Cabin/Bungalow	6.955e-01	7.330e-01	0.949	0.342706
room_type_property_group_interactionEntire home/apt.Guest Accommodation	8.468e-01	1.260e-01	6.722	1.79e-11 ***
room_type_property_group_interactionPrivate room.Guest Accommodation	1.171e+00	3.218e-01	3.640	0.000273 ***
room_type_property_group_interactionHotel room.Hotel/Hostel	-1.302e+01	1.549e+02	-0.084	0.933015
room_type_property_group_interactionPrivate room.Hotel/Hostel	-1.899e+00	5.940e-01	-3.197	0.001388 **
room_type_property_group_interactionShared room.Hotel/Hostel	-1.208e+01	3.845e+02	-0.031	0.974946
room_type_property_group_interactionEntire home/apt.House/Townhouse/Villa	1.864e-01	8.572e-02	2.174	0.029670 *
room_type_property_group_interactionPrivate room.House/Townhouse/Villa	9.779e-01	1.258e-01	7.771	7.81e-15 ***
room_type_property_group_interactionEntire home/apt.Unique Stays (Other)	4.374e-01	4.292e-01	1.019	0.308121
room_type_property_group_interactionPrivate room.Unique Stays (Other)	1.256e+00	1.443e-01	8.708	< 2e-16 ***
room_type_property_group_interactionShared room.Unique Stays (Other)	2.361e-01	6.439e-01	0.367	0.713816

```

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 7145.5 on 5757 degrees of freedom
Residual deviance: 6373.4 on 5734 degrees of freedom
AIC: 6421.4

```

Number of Fisher Scoring iterations: 13

Final Logistic regression summary

```

Call:
glm(formula = HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
     high_availability + longitude + instant_bookable + region +
     price_availability_interaction + room_type_property_group_interaction,
     family = binomial, data = train)

```

```

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.7912   -0.8864   -0.6180    1.1273    3.0352

```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.298e+02	6.333e+01	-2.050	0.040388 *
host_is_superhost	5.361e-01	6.338e-02	8.458	< 2e-16 ***
log_price	7.267e-01	8.835e-02	8.225	< 2e-16 ***
calculated_host_listings_count	-1.805e-02	1.793e-03	-10.066	< 2e-16 ***
high_availability1	-3.465e-01	8.245e-02	-4.203	2.63e-05 ***
longitude	8.272e-01	4.200e-01	1.969	0.048916 *
instant_bookable	-4.813e-01	7.320e-02	-6.575	4.87e-11 ***
regionEast	-2.232e-02	1.431e-01	-0.156	0.876046
regionNorth	2.100e-01	1.483e-01	1.416	0.156845
regionSouth	3.523e-01	1.579e-01	2.231	0.025711 *
price_availability_interaction	-7.221e-06	5.325e-06	-1.356	0.175117
room_type_property_group_interactionHotel room.Apartment/Condo	2.274e-01	8.751e-01	0.260	0.794928
room_type_property_group_interactionPrivate room.Apartment/Condo	1.027e+00	3.444e-01	2.982	0.002866 **
room_type_property_group_interactionEntire home/apt.Cottage/Cabin/Bungalow	7.475e-01	2.420e-01	3.089	0.002009 **
room_type_property_group_interactionPrivate room.Cottage/Cabin/Bungalow	7.438e-01	7.319e-01	1.016	0.309475
room_type_property_group_interactionEntire home/apt.Guest Accommodation	8.555e-01	1.261e-01	6.782	1.19e-11 ***
room_type_property_group_interactionPrivate room.Guest Accommodation	1.192e+00	3.221e-01	3.700	0.000215 ***
room_type_property_group_interactionHotel room.Hotel/Hostel	-1.301e+01	1.550e+02	-0.084	0.933090
room_type_property_group_interactionPrivate room.Hotel/Hostel	-1.887e+00	5.939e-01	-3.177	0.001486 **
room_type_property_group_interactionShared room.Hotel/Hostel	-1.204e+01	3.844e+02	-0.031	0.975005
room_type_property_group_interactionEntire home/apt.House/Townhouse/Villa	1.900e-01	8.578e-02	2.215	0.026727 *
room_type_property_group_interactionPrivate room.House/Townhouse/Villa	1.008e+00	1.278e-01	7.887	3.10e-15 ***
room_type_property_group_interactionEntire home/apt.Unique Stays (Other)	4.510e-01	4.288e-01	1.052	0.292879
room_type_property_group_interactionPrivate room.Unique Stays (Other)	1.279e+00	1.452e-01	8.803	< 2e-16 ***
room_type_property_group_interactionShared room.Unique Stays (Other)	2.886e-01	6.453e-01	0.447	0.654683

```

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 7145.5 on 5757 degrees of freedom
Residual deviance: 6371.5 on 5733 degrees of freedom
AIC: 6421.5

```

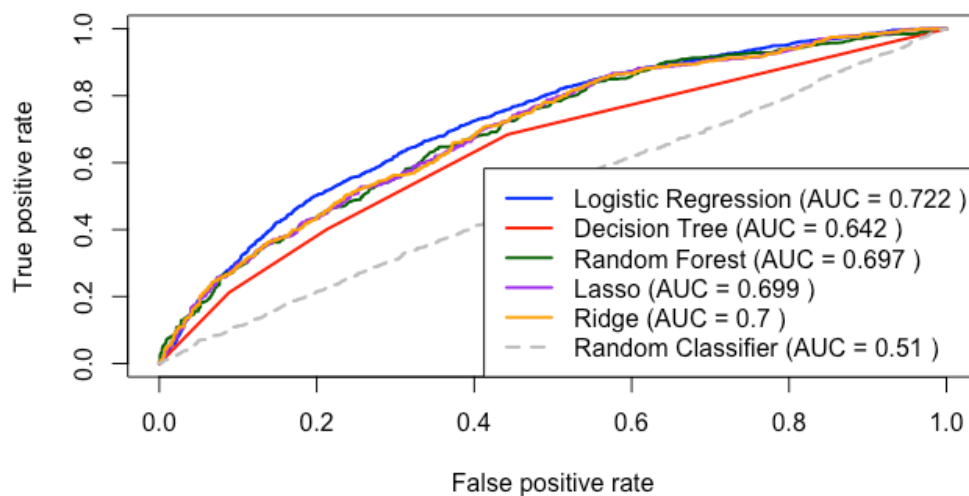
Number of Fisher Scoring iterations: 13

2.5 Model Comparison

Model	Accuracy	Precision	Recall	F1-Score	AUC	Interpretability
Logistic Regression	0.713	0.601	0.601	0.601	0.722	High
Decision Tree	0.694	0.519	0.519	0.519	0.642	Very High
Random Forest	0.697	0.518	0.368	0.430	0.697	Lower (black box)
Lasso Regression	0.712	0.583	0.258	0.358	0.699	High (sparse)
Ridge Regression	0.717	0.609	0.249	0.353	0.700	High (shrinkage)

ROC curves and AUC scores for all models

ROC Curve Comparison



2.6 Model Summary (Logistic Regression)

Model Fit

Metric	Value
Null Deviance	7145.5
Residual Deviance	6371.5
AIC	6421.5

5.4 Model comparison

Model	Val MSE	Val RMSE (\$)	Test MSE	Test RMSE	Comments
XGBoost	11,621.14	107.80	11,602.35	107.71	Best predictive performance
GBM	12,125.03	110.11	12,487.32	111.75	Strong performance, slightly behind XGB
Random Forest	12,180.51	110.37	12,829.36	113.27	Competitive model
GLM	14,763.91	121.51	15,917.98	126.17	Simple interpretable baseline
Lasso Reg.	14,780.49	121.58	15,921.67	126.18	Regularized, but weaker prediction

Code

```

```{r}
Load necessary libraries
library(randomForest)
library(MASS)
library(readxl)
library(glmnet)
library(caret)
library(tidyverse)
library(ggplot2)
library(corrplot)
library(dplyr)
library(psych)
library(ROCR)
library(rpart)
library(splines)
library(xgboost)
library(Matrix)
library(gbm)
library(writexl)
```

## Q1

```{r}
df <- read_excel("/Users/audreychang/Desktop/ACTL_ML/Textbook_R/HW/Assignment2/AirbnbSydneyV2.xlsx")
```

### Create avg_score from the 7 review score variables

```{r}
df$avg_score <- rowMeans(df[, c("review_scores_rating", "review_scores_accuracy",
 "review_scores_cleanliness", "review_scores_checkin",
 "review_scores_communication", "review_scores_location",

```

```

 "review_scores_value")], na.rm = TRUE)
 ...

  ```{r}
  # Create binary HAS variable: 1 if avg_score > 4.9, else 0
  df$HAS <- as.numeric(df$avg_score > 4.9)

  ...

  ```{r}
 prop_HAS <- mean(df$HAS)
 cat("Proportion of listings with HAS=1:", prop_HAS, "\n")
 ...

Q2

Data Preprocessing and Feature Engineering

Clean and transform data

  ```{r}
  # Convert character indicators to numeric
  df <- df %>%
    mutate(
      host_is_superhost = ifelse(host_is_superhost == "t", 1, 0),
      host_identity_verified = ifelse(host_identity_verified == "t", 1, 0),
      host_has_profile_pic = ifelse(host_has_profile_pic == "t", 1, 0),
      instant_bookable = ifelse(instant_bookable == "t", 1, 0),
      host_experience_years = 2025 - lubridate::year(host_since),
      host_listings_count = as.numeric(host_listings_count),
      host_total_listings_count = as.numeric(host_total_listings_count),
      host_response_rate = as.numeric(host_response_rate),
      host_acceptance_rate = as.numeric(host_acceptance_rate)
    )

  ...

  ```{r}
 # Handle missing values
 # Replace "N/A" strings with actual NA values
 na_string_cols <- sapply(df, function(x) {
 if (is.character(x) || is.factor(x)) {
 any(x == "N/A", na.rm = TRUE)
 } else {
 FALSE
 }
 })
 cols_to_fix <- names(na_string_cols[na_string_cols == TRUE])

```

```

df[cols_to_fix] <- lapply(df[cols_to_fix], function(x) na_if(x, "N/A"))

Check NA values in each column
na_count <- sapply(df, function(x) sum(is.na(x)))
na_count[na_count > 0] # Show columns with missing values

Calculate proportion of missing values
na_proportion <- sapply(df, function(x) mean(is.na(x)))

Remove columns not needed
df <- df[, !(names(df) %in% c("bathrooms_text", "host_since"))]

Replace missing values in categorical variables
df$host_response_time[is.na(df$host_response_time)] <- "Unknown"

Impute host_response_rate with median
median_response_rate <- median(df$host_response_rate, na.rm = TRUE)
df$host_response_rate[is.na(df$host_response_rate)] <- median_response_rate

Impute host_acceptance_rate with median
median_acceptance_rate <- median(df$host_acceptance_rate, na.rm = TRUE)
df$host_acceptance_rate[is.na(df$host_acceptance_rate)] <- median_acceptance_rate

Impute host_is_superhost with 0 (Not Superhost) most common type
Since it's already numeric (0 = not superhost, 1 = superhost)
df$host_is_superhost[is.na(df$host_is_superhost)] <- 0
...

Feature Engineering: Create new variables

```{r}
# Create region categories
east <- c("Randwick", "Waverley", "Woollahra", "Sydney", "Marrickville",
        "Canada Bay", "Botany Bay", "Leichhardt", "Ashfield", "Burwood")
west <- c("Penrith", "Blacktown", "Fairfield", "Liverpool", "Campbelltown",
        "Auburn", "Bankstown", "Holroyd", "Strathfield", "Parramatta", "Canterbury")
north <- c("Ku-Ring-Gai", "North Sydney", "Willoughby", "Lane Cove", "Hornsby",
        "Ryde", "Mosman", "Hunters Hill", "The Hills Shire", "Warringah", "Manly", "Pittwater")
south <- c("Sutherland Shire", "Rockdale", "Hurstville", "Bankstown", "Camden", "City Of Kogarah")

# Create region classification
df$region <- case_when(
  df$neighbourhood %in% east ~ "East",
  df$neighbourhood %in% west ~ "West",
  df$neighbourhood %in% north ~ "North",
  df$neighbourhood %in% south ~ "South",
  TRUE ~ "Other"
)

```

```

# Categorize property types into broad groups
df$property_group <- ifelse(df$property_type %in% c(
  "Entire rental unit", "Entire condo", "Entire serviced apartment",
  "Entire loft", "Room in serviced apartment", "Room in aparthotel",
  "Private room in serviced apartment", "Private room in condo"),
  "Apartment/Condo",
  ifelse(df$property_type %in% c(
    "Entire home", "Entire vacation home", "Entire villa",
    "Entire townhouse", "Private room in home", "Private room in townhouse",
    "Private room in villa"),
    "House/Townhouse/Villa",
    ifelse(df$property_type %in% c(
      "Entire guesthouse", "Entire guest suite",
      "Private room in guesthouse", "Private room in guest suite",
      "Private room in bed and breakfast"),
      "Guest Accommodation",
      ifelse(df$property_type %in% c(
        "Entire cottage", "Entire cabin", "Entire bungalow",
        "Private room in cabin", "Private room in cottage",
        "Private room in bungalow", "Private room in chalet", "Tiny home"),
        "Cottage/Cabin/Bungalow",
        ifelse(df$property_type %in% c(
          "Room in hotel", "Room in boutique hotel", "Shared room in boutique hotel",
          "Shared room in hostel", "Private room in hostel",
          "Shared room in hotel", "Shared room in condo"),
          "Hotel/Hostel",
          "Unique Stays (Other)"
        )
      )
    )
  )
))

```

```

# Create high_availability indicator
df$high_availability <- ifelse(df$availability_60 > median(df$availability_60), 1, 0)

```

```

# Drop original variables after categorization
df <- df[, !(names(df) %in% c("neighbourhood", "property_type"))]

```

```

# Convert categorical variables to factors
df$room_type <- as.factor(df$room_type)
df$property_group <- as.factor(df$property_group)
df$region <- as.factor(df$region)
df$high_availability <- as.factor(df$high_availability)

```

```
```
```

### ### Exploratory Data Analysis

```
```{r}
```

```
# Create a clean dataset excluding review-based variables
```

```
df_clean <- df
```

```
columns_to_exclude <- c("review_scores_rating",  
                        "review_scores_accuracy",  
                        "review_scores_cleanliness",  
                        "review_scores_checkin",  
                        "review_scores_communication",  
                        "review_scores_location",  
                        "review_scores_value",  
                        "avg_score")
```

```
df_clean <- df_clean[, !(names(df_clean) %in% columns_to_exclude)]
```

```
# Ensure HAS is numeric for correlation analysis
```

```
df_clean$HAS <- as.numeric(df_clean$HAS)
```

```
# Calculate correlation with HAS for numeric variables
```

```
numeric_vars <- df_clean[, sapply(df_clean, is.numeric)]
```

```
numeric_vars <- numeric_vars[, names(numeric_vars) != "HAS"]
```

```
cor_results <- sapply(numeric_vars, function(x) {  
  if (is.numeric(x)) {  
    cor(df_clean$HAS, x, use = "complete.obs")  
  } else {  
    NA  
  }  
})
```

```
# Remove NA results
```

```
cor_results <- cor_results[!is.na(cor_results)]
```

```
# Create correlation summary dataframe
```

```
cor_df <- data.frame(  
  Variable = names(cor_results),  
  CorrelationWithHAS = cor_results  
) %>% arrange(desc(abs(CorrelationWithHAS)))
```

```
print(cor_df)
```

```
# Visualize correlations
```

```
ggplot(cor_df, aes(x = reorder(Variable, CorrelationWithHAS), y = CorrelationWithHAS)) +  
  geom_col(fill = "steelblue") +  
  geom_text(aes(label = round(CorrelationWithHAS, 3)),  
    hjust = ifelse(cor_df$CorrelationWithHAS >= 0, -0.1, 1.1),
```

```

      size = 3.5) +
coord_flip() +
labs(title = "Point-Biserial Correlation with HAS",
      x = "Variable",
      y = "Correlation") +
theme(axis.text.x = element_text(face = "bold"),
      axis.text.y = element_text(face = "bold"),
      plot.title = element_text(hjust = 0.5, face = "bold"))
...

```{r}
Convert HAS to factor for visualizations and modeling
df$HAS <- factor(df$HAS)

Visualize categorical variables
categorical_vars <- c("region", "property_group", "host_response_time", "room_type", "high_availability")

Loop through each categorical variable and create a visualization
for (var in categorical_vars) {
 # Create a summary table
 summary_table <- df %>%
 group_by(HAS, !!sym(var)) %>%
 summarise(count = n(), .groups = "drop")

 # Plot the data
 p <- ggplot(summary_table, aes_string(x = var, y = "count", fill = "HAS")) +
 geom_bar(stat = "identity", position = "stack") +
 labs(title = paste("Distribution of", var, "by HAS"),
 x = var, y = "Count") +
 theme_minimal() +
 scale_fill_manual(values = c("0" = "blue", "1" = "red")) # Match fill to factor levels

 # Rotate x-axis labels for crowded variables
 if (var %in% c("property_group", "region", "host_response_time", "high_availability")) {
 p <- p + theme(axis.text.x = element_text(angle = 45, hjust = 1))
 }

 print(p)
}
...

Feature Transformation and Interaction based on EDA

```{r}
# create transformations and interactions based on insights from EDA

# Set reference levels for categorical variables

```

```

# First check what levels exist
region_levels <- levels(df$region)
print(region_levels)

# Relevel with appropriate error handling
if("West" %in% region_levels) {
  df$region <- relevel(df$region, ref = "West")
} else {
  cat("Warning: 'West' not found in region levels. Available levels:", paste(region_levels, collapse=", "), "\n")
  # Use first level if "West" isn't available
  df$region <- relevel(df$region, ref = region_levels[1])
}

df$property_group <- relevel(df$property_group, ref = "Apartment/Condo")
df$room_type <- relevel(df$room_type, ref = "Entire home/apt")

# log transformation to handle skewness
df$log_price <- log(df$price + 1)

# Create interactions based on EDA insights
df$host_exp_superhost_interaction <- df$host_experience_years * df$host_is_superhost
df$price_availability_interaction <- df$price * df$availability_60
df$room_type_property_group_interaction <- interaction(df$room_type, df$property_group)
...

## Q2,3,4

### Model Building

```{r}
First make sure HAS is a factor for classification
df$HAS <- factor(df$HAS)

Split data into training and test sets
set.seed(123)
split <- sample(1:nrow(df), 0.75 * nrow(df))
train <- df[split,]
test <- df[-split,]

Handle missing values in training data
For numeric columns, replace NA with median
train <- train %>%
 mutate(across(where(is.numeric), ~ifelse(is.na(.), median(., na.rm = TRUE), .)))

For categorical variables, replace NA with most common category
for(col in names(train)[sapply(train, is.factor)]) {

```

```

if(any(is.na(train[[col]]))) {
 most_common <- names(sort(table(train[[col]]), decreasing = TRUE))[1]
 train[[col]][is.na(train[[col]])] <- most_common
}
}

Apply same transformations to test set
test <- test %>%
 mutate(across(where(is.numeric), ~ifelse(is.na(.), median(., na.rm = TRUE), .)))

for(col in names(test)[sapply(test, is.factor)]) {
 if(any(is.na(test[[col]]))) {
 most_common <- names(sort(table(train[[col]]), decreasing = TRUE))[1] # Use distribution from training data
 test[[col]][is.na(test[[col]])] <- most_common
 }
}

Extract response variables
y_train <- train$HAS
y_test <- test$HAS
```

```

Train-Test Split

```

```{r}
Split
set.seed(123)
split <- sample(1:nrow(df), 0.75 * nrow(df))
train <- df[split,]
test <- df[-split,]

If filtering, do it here
train <- na.omit(train)
test <- na.omit(test)

Then extract y
y_train <- train$HAS
y_test <- test$HAS
```

```

Logistic Regression (Baseline)

1. Stepwise Selection

```

```{r}
Start with full model
logit_full <- glm(HAS ~ host_is_superhost + log_price + calculated_host_listings_count +

```



```

 high_availability + longitude + latitude + instant_bookable +
 region + price_availability_interaction + room_type_property_group_interaction,
 family = binomial, data = train)

```

```

Run stepwise selection

```

```

model_stepwise <- stepAIC(logit_full, direction = "both")
summary(model_stepwise)

```

```

...

```

```

```{r}

```

```

# Final Logistic regression model

```

```

logit_revised <- glm(HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
    high_availability + longitude + instant_bookable +
    region + price_availability_interaction +
    room_type_property_group_interaction,
    family = binomial, data = train)

```

```

# View the summary of the updated model

```

```

summary(logit_revised)

```

```

...

```

```

```{r}

```

```

Get predicted probabilities for the logistic model

```

```

pred_probs <- predict(logit_revised, type = "response")

```

```

Convert probabilities to binary labels using a threshold of 0.5

```

```

pred_labels <- ifelse(pred_probs > 0.5, 1, 0)

```

```

Actual labels (assuming your dependent variable is 'HAS')

```

```

true_labels <- train$HAS

```

```

Confusion Matrix

```

```

cm <- table(Predicted = pred_labels, Actual = true_labels)
print(cm)

```

```

Accuracy

```

```

accuracy <- sum(diag(cm)) / sum(cm)
cat("Accuracy:", accuracy, "\n")

```

```

Precision

```

```

precision <- cm[2, 2] / sum(cm[2,]) # TP / (TP + FP)
cat("Precision:", precision, "\n")

```

```

Recall (Sensitivity)

```

```

recall <- cm[2, 2] / sum(cm[2,]) # TP / (TP + FN)
cat("Recall:", recall, "\n")

```

```

F1-score
f1 <- 2 * (precision * recall) / (precision + recall)
cat("F1-score:", f1, "\n")

ROC Curve & AUC
pred_logit <- prediction(pred_probs, true_labels)
perf <- performance(pred_logit, measure = "tpr", x.measure = "fpr")
plot(perf, main = "ROC Curve", col = "blue", lwd = 2)
roc_auc <- performance(pred_logit, measure = "auc")@y.values[[1]]
abline(a = 0, b = 1, col = "gray", lty = 2)
cat("AUC:", roc_auc, "\n")
``

```

## Decision Tree

```

``{r}
Train the Decision Tree model
dt_model <- rpart(HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
 high_availability + longitude + instant_bookable +
 region + price_availability_interaction +
 room_type_property_group_interaction,
 data = train, method = "class")

Make predictions on the test set
pred_probs_dt <- predict(dt_model, test, type = "prob")[, 2] # Get probabilities for the positive class
pred_labels_dt <- ifelse(pred_probs_dt > 0.5, 1, 0) # Convert probabilities to binary labels

Actual labels for the test set
true_labels_dt <- test$HAS

Evaluate the model
Confusion Matrix
cm_dt <- table(Predicted = pred_labels_dt, Actual = true_labels_dt)
print(cm_dt)

Accuracy
accuracy_dt <- sum(diag(cm_dt)) / sum(cm_dt)
cat("Accuracy:", accuracy_dt, "\n")

Precision
precision_dt <- cm_dt[2, 2] / sum(cm_dt[2,]) # TP / (TP + FP)
cat("Precision:", precision_dt, "\n")

Recall (Sensitivity)
recall_dt <- cm_dt[2, 2] / sum(cm_dt[2,]) # TP / (TP + FN)
cat("Recall:", recall_dt, "\n")

```

```

F1-score
f1_dt <- 2 * (precision_dt * recall_dt) / (precision_dt + recall_dt)
cat("F1-score:", f1_dt, "\n")

ROC Curve & AUC
pred_dt <- prediction(pred_probs_dt, true_labels_dt)
perf_dt <- performance(pred_dt, measure = "tpr", x.measure = "fpr")
plot(perf_dt, main = "ROC Curve - Decision Tree", col = "red", lwd = 2)
roc_auc_dt <- performance(pred_dt, measure = "auc")@y.values[[1]]
abline(a = 0, b = 1, col = "gray", lty = 2)
cat("AUC:", roc_auc_dt, "\n")
```

#### Random Forest Model

```{r}
Train Random Forest
set.seed(123)
rf_model <- randomForest(HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
 high_availability + longitude + instant_bookable + region +
 price_availability_interaction + room_type_property_group_interaction,
 data = train, importance = TRUE, ntree = 500)

Predict probabilities and classes on the test set
rf_probs <- predict(rf_model, newdata = test, type = "prob")[, 2] # probabilities for class 1
rf_preds <- ifelse(rf_probs > 0.5, 1, 0)

Convert numeric predictions to factors for confusionMatrix function
rf_preds_factor <- factor(rf_preds, levels = c(0, 1))
test_HAS_factor <- factor(as.numeric(as.character(test$HAS)), levels = c(0, 1))

Confusion Matrix & Basic Metrics
conf_rf <- confusionMatrix(rf_preds_factor, test_HAS_factor, positive = "1")
print(conf_rf)

Extract evaluation metrics
accuracy_rf <- conf_rf$overall["Accuracy"]
precision_rf <- conf_rf$byClass["Pos Pred Value"]
recall_rf <- conf_rf$byClass["Sensitivity"]
f1_rf <- conf_rf$byClass["F1"]

cat("\n--- Evaluation Metrics: Random Forest ---\n")
cat("Accuracy :", round(accuracy_rf, 4), "\n")
cat("Precision:", round(precision_rf, 4), "\n")
cat("Recall :", round(recall_rf, 4), "\n")
cat("F1-score :", round(f1_rf, 4), "\n")

ROC Curve & AUC

```

```

pred_rf <- prediction(rf_probs, as.numeric(as.character(test$HAS)))
perf_rf <- performance(pred_rf, "tpr", "fpr")
roc_auc_rf <- performance(pred_rf, "auc")@y.values[[1]]

Plot ROC
plot(perf_rf, main = "ROC Curve - Random Forest", col = "darkgreen", lwd = 2)
abline(a = 0, b = 1, col = "gray", lty = 2)
cat("AUC :", round(roc_auc_rf, 4), "\n")

Variable importance plot
varImpPlot(rf_model, main = "Variable Importance - Random Forest")
```

```

Regularized Logistic Regression (Lasso/Ridge)

Prepare Data

```

```{r}
--- Prepare data ---
Build model matrix from predictors
x_train <- model.matrix(HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
 high_availability + longitude + instant_bookable +
 region + price_availability_interaction +
 room_type_property_group_interaction, data = train)[, -1] # drop intercept

x_test <- model.matrix(HAS ~ host_is_superhost + log_price + calculated_host_listings_count +
 high_availability + longitude + instant_bookable +
 region + price_availability_interaction +
 room_type_property_group_interaction, data = test)[, -1] # drop intercept

```

```

Lasso Logistic Regression (L1)

```

```{r}
Lasso (alpha = 1)
lasso_model <- cv.glmnet(x_train, y_train, alpha = 1, family = "binomial", type.measure = "class")

Predict on test set
lasso_probs <- predict(lasso_model, newx = x_test, s = "lambda.min", type = "response")
lasso_preds <- ifelse(lasso_probs > 0.5, 1, 0)

Evaluation
conf_lasso <- confusionMatrix(as.factor(lasso_preds), as.factor(y_test), positive = "1")
print(conf_lasso)

Metrics
accuracy_lasso <- conf_lasso$overall["Accuracy"]

```

```
precision_lasso <- conf_lasso$byClass["Precision"]
recall_lasso <- conf_lasso$byClass["Recall"]
f1_lasso <- conf_lasso$byClass["F1"]
```

```
cat("\n--- Lasso Logistic Regression ---\n")
cat("Accuracy :", round(accuracy_lasso, 4), "\n")
cat("Precision:", round(precision_lasso, 4), "\n")
cat("Recall :", round(recall_lasso, 4), "\n")
cat("F1-score :", round(f1_lasso, 4), "\n")
```

```
```
```

```
```{r}
Check levels of the target variable
levels(y_train)
levels(y_test)
```
```

Ridge Logistic Regression (L2)

```
```{r}
Ridge (alpha = 0)
ridge_model <- cv.glmnet(x_train, y_train, alpha = 0, family = "binomial", type.measure = "class")
```

```
Predict on test set
ridge_probs <- predict(ridge_model, newx = x_test, s = "lambda.min", type = "response")
ridge_preds <- ifelse(ridge_probs > 0.5, 1, 0)
```

```
Evaluation
conf_ridge <- confusionMatrix(as.factor(ridge_preds), as.factor(y_test), positive = "1")
print(conf_ridge)
```

```
Metrics
accuracy_ridge <- conf_ridge$overall["Accuracy"]
precision_ridge <- conf_ridge$byClass["Precision"]
recall_ridge <- conf_ridge$byClass["Recall"]
f1_ridge <- conf_ridge$byClass["F1"]
```

```
cat("\n--- Ridge Logistic Regression ---\n")
cat("Accuracy :", round(accuracy_ridge, 4), "\n")
cat("Precision:", round(precision_ridge, 4), "\n")
cat("Recall :", round(recall_ridge, 4), "\n")
cat("F1-score :", round(f1_ridge, 4), "\n")
```

```
```
```

ROC Curve & AUC for both

```

```{r}
Lasso
pred_lasso <- prediction(lasso_probs, y_test)
perf_lasso <- performance(pred_lasso, "tpr", "fpr")
auc_lasso <- performance(pred_lasso, "auc")@y.values[[1]]

Ridge
pred_ridge <- prediction(ridge_probs, y_test)
perf_ridge <- performance(pred_ridge, "tpr", "fpr")
auc_ridge <- performance(pred_ridge, "auc")@y.values[[1]]

Plot ROC
plot(perf_lasso, main = "ROC Curve - Regularized Logistic Regression", col = "purple", lwd = 2)
plot(perf_ridge, add = TRUE, col = "orange", lwd = 2)
legend("bottomright",
 legend = c(
 paste("Lasso (AUC =", round(auc_lasso, 3), ")"),
 paste("Ridge (AUC =", round(auc_ridge, 3), ")")
),
 col = c("purple", "orange"), lwd = 2)

cat("Lasso AUC:", auc_lasso, "\n")
cat("Ridge AUC:", auc_ridge, "\n")

```

```

```

```

ROC comparison

```

```{r}
--- Logistic Regression ROC ---
pred_logit <- prediction(pred_probs, true_labels)
perf_logit <- performance(pred_logit, measure = "tpr", x.measure = "fpr")
roc_auc_logit <- performance(pred_logit, measure = "auc")@y.values[[1]]

--- Decision Tree ROC ---
pred_dt <- prediction(pred_probs_dt, true_labels_dt)
perf_dt <- performance(pred_dt, measure = "tpr", x.measure = "fpr")
roc_auc_dt <- performance(pred_dt, measure = "auc")@y.values[[1]]

--- Random Forest ROC ---
pred_rf <- prediction(rf_probs, test$HAS)
perf_rf <- performance(pred_rf, measure = "tpr", x.measure = "fpr")
roc_auc_rf <- performance(pred_rf, measure = "auc")@y.values[[1]]

--- Lasso ROC ---
pred_lasso <- prediction(lasso_probs, y_test)
perf_lasso <- performance(pred_lasso, "tpr", "fpr")
roc_auc_lasso <- performance(pred_lasso, "auc")@y.values[[1]]

```

```

--- Ridge ROC ---
pred_ridge <- prediction(ridge_probs, y_test)
perf_ridge <- performance(pred_ridge, "tpr", "fpr")
roc_auc_ridge <- performance(pred_ridge, "auc")@y.values[[1]]

--- Random Classifier (baseline) ---
set.seed(42)
random_probs <- runif(length(true_labels)) # generate random probabilities
pred_random <- prediction(random_probs, true_labels)
perf_random <- performance(pred_random, "tpr", "fpr")
roc_auc_random <- performance(pred_random, "auc")@y.values[[1]]

--- Plot ROC curves together ---
plot(perf_logit, main = "ROC Curve Comparison", col = "blue", lwd = 2)
plot(perf_dt, add = TRUE, col = "red", lwd = 2)
plot(perf_rf, add = TRUE, col = "darkgreen", lwd = 2)
plot(perf_lasso, add = TRUE, col = "purple", lwd = 2)
plot(perf_ridge, add = TRUE, col = "orange", lwd = 2)
plot(perf_random, add = TRUE, col = "gray", lwd = 2, lty = 2)

--- Add the legend ---
legend("bottomright",
 legend = c(
 paste("Logistic Regression (AUC =", round(roc_auc_logit, 3), ")"),
 paste("Decision Tree (AUC =", round(roc_auc_dt, 3), ")"),
 paste("Random Forest (AUC =", round(roc_auc_rf, 3), ")"),
 paste("Lasso (AUC =", round(roc_auc_lasso, 3), ")"),
 paste("Ridge (AUC =", round(roc_auc_ridge, 3), ")"),
 paste("Random Classifier (AUC =", round(roc_auc_random, 3), ")")
),
 col = c("blue", "red", "darkgreen", "purple", "orange", "gray"),
 lty = c(1, 1, 1, 1, 1, 2),
 lwd = 2)

--- AUC output ---
cat("Logistic Regression AUC :", roc_auc_logit, "\n")
cat("Decision Tree AUC :", roc_auc_dt, "\n")
cat("Random Forest AUC :", roc_auc_rf, "\n")
cat("Lasso Logistic AUC :", roc_auc_lasso, "\n")
cat("Ridge Logistic AUC :", roc_auc_ridge, "\n")
cat("Random Classifier AUC :", roc_auc_random, "\n")

```

```

...

```

Q3

```

...{r}

```

```

--- Residual and Deviance Diagnostics for Final Logistic Regression Model ---

Pearson residuals
pearson_resid <- residuals(logit_revised, type = "pearson")

Deviance residuals
deviance_resid <- residuals(logit_revised, type = "deviance")

Fitted values (predicted probabilities)
fitted_probs <- fitted(logit_revised)

--- Plot: Deviance Residuals vs Fitted Probabilities ---

ggplot(data = data.frame(Fitted = fitted_probs, Deviance = deviance_resid), aes(x = Fitted, y = Deviance)) +
 geom_point(alpha = 0.5, color = "steelblue") +
 geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
 labs(
 title = "Deviance Residuals vs Fitted Probabilities",
 x = "Fitted Probabilities",
 y = "Deviance Residuals"
) +
 theme_minimal()

...

```{r}
# --- Influence Diagnostics ---
library(ggplot2)

# Cook's Distance
cooks_d <- cooks.distance(logit_revised)

# Leverage (hat) values
hat_values <- hatvalues(logit_revised)

# Combine diagnostics
influence_df <- data.frame(
  obs = 1:length(cooks_d),
  cook = cooks_d,
  hat = hat_values,
  deviance = residuals(logit_revised, type = "deviance")
)

# --- Plot Cook's Distance ---
ggplot(influence_df, aes(x = obs, y = cook)) +
  geom_bar(stat = "identity", fill = "firebrick", alpha = 0.7) +
  geom_hline(yintercept = 4 / nrow(influence_df), linetype = "dashed", color = "black") +
  labs(

```



```

    title = "Cook's Distance for Logistic Regression",
    x = "Observation Index",
    y = "Cook's Distance"
  ) +
  theme_minimal()

...

## Q5

```{r}
Preprocessing and Feature Engineering

Drop review-related variables Keep avg_score to be the proxy of all review-score-related variables
review_vars <- c(
 "review_scores_rating",
 "review_scores_accuracy",
 "review_scores_cleanliness",
 "review_scores_checkin",
 "review_scores_communication",
 "review_scores_location",
 "review_scores_value"
)

Create df_clean2 by removing review variables
df_clean2 <- df[, !(names(df) %in% review_vars)]

Remove other unusable columns
df_clean2 <- df_clean2[, !(names(df_clean2) %in% c(
 "price", # Remove original price as I'll use log_price
 "price_availability_interaction",
 "high_availability"
)))]

Feature Engineering: Calculate days between reviews
This might introduce data leakage for new listings - only use if available at prediction time
df_clean2$days_between_reviews <- as.numeric(difftime(
 as.Date(df_clean2$last_review),
 as.Date(df_clean2$first_review),
 units = "days"
))

Remove original date columns after feature extraction
df_clean2 <- df_clean2[, !(names(df_clean2) %in% c("first_review", "last_review"))]

Handle missing values
Replace "N/A" strings with actual NA

```

```

na_string_cols <- sapply(df_clean2, function(x) {
 if (is.character(x) || is.factor(x)) {
 any(x == "N/A", na.rm = TRUE)
 } else {
 FALSE
 }
})

cols_to_fix <- names(na_string_cols[na_string_cols == TRUE])
df_clean2[cols_to_fix] <- lapply(df_clean2[cols_to_fix], function(x) na_if(x, "N/A"))

Check for remaining NA values
na_count <- sapply(df_clean2, function(x) sum(is.na(x)))
na_count[na_count > 0] # Show only columns with missing values
```



```

```{r}
# Select all numeric predictors for correlation analysis
numeric_vars <- names(df_clean2)[sapply(df_clean2, is.numeric)]
df_numeric <- df_clean2[numeric_vars]

# Calculate correlation with log_price
cor_with_price <- sapply(numeric_vars, function(x) {
  if(x != "log_price") {
    cor(df_clean2[[x]], df_clean2$log_price, use = "complete.obs")
  } else {
    NA
  }
})

# Create and sort correlation dataframe
cor_df <- data.frame(
  Predictor = names(cor_with_price),
  Correlation = cor_with_price
) %>%
  filter(!is.na(Correlation)) %>% # Remove NA values
  arrange(desc(abs(Correlation))) # Sort by absolute correlation for both pos/neg correlations

# Print the sorted correlations
print(cor_df)

# Plot correlation
ggplot(cor_df, aes(x = reorder(Predictor, Correlation), y = Correlation)) +
  geom_bar(stat = "identity", fill = "skyblue", color = "black") +
  coord_flip() +
  theme_minimal(base_size = 15) +
  labs(

```


```

```

 title = "Correlation of Numeric Predictors with Airbnb Log_Price",
 x = "Predictor",
 y = "Correlation with Log_Price"
) +
 theme(
 plot.title = element_text(size = 18, face = "bold", hjust = 0.5),
 axis.title = element_text(size = 14, face = "bold"),
 axis.text = element_text(size = 12, face = "bold")
) +
 geom_text(
 aes(label = round(Correlation, 2)),
 hjust = ifelse(cor_df$Correlation > 0, -0.2, 1.2),
 size = 4,
 fontface = "bold"
)
}

```

## Prepossessing data

\\*\\*PCA on correlated subset\\*\\*:

```

```{r}
# Select highly correlated "capacity" variables
capacity_vars <- df_clean2 %>% select(accommodates, beds, bedrooms, bathrooms)

# Perform PCA
pca_capacity <- prcomp(capacity_vars, center = TRUE, scale. = TRUE)

# Add first principal component back to data
df_clean2$PC_capacity <- pca_capacity$x[, 1]
#Drop original correlated variables
df_clean2 <- df_clean2 %>% select(-accommodates, -beds, -bedrooms, -bathrooms)

```

```{r}
colnames(df_clean2)
#Remove rows with missing values
df_clean2 <- na.omit(df_clean2)

```

```{r}
# Fit initial Random Forest model to get variable importance
rf_model <- randomForest(log_price ~ ., data = df_clean2, ntree = 500, importance = TRUE)

# Extract variable importance
importance_df <- as.data.frame(importance(rf_model))
importance_df$Variable <- rownames(importance_df)

```

```

# Sort by %IncMSE (most important at the top)
importance_sorted <- importance_df %>%
  arrange(desc(`%IncMSE`))

# Print the top variables by importance
print(importance_sorted)

# Plot variable importance
varImpPlot(rf_model, main = "Variable Importance (%IncMSE)")

...

```{r}
#Try several cutoffs:
mse_results <- c()
for (n in c(10, 15, 20)) {
 top_n_vars <- importance_sorted$Variable[1:n]
 df_sub <- df_clean2 %>% select(all_of(top_n_vars), log_price)

 # Fit model (Random Forest)
 rf_temp <- randomForest(log_price ~ ., data = df_sub, ntree = 500)
 pred_temp <- predict(rf_temp)
 mse <- mean((df_sub$log_price - pred_temp)^2)
 mse_results <- c(mse_results, mse)
}
print(mse_results)

...

Create the reduced dataset (`df_reduced`) using the selected top 20 predictors.

```{r}
# Extract top 20 predictor variable names from importance
top_20_vars <- importance_sorted$Variable[1:20]

# Filter top 20 variables to include only those that exist in df_clean2
top_20_vars <- top_20_vars[top_20_vars %in% names(df_clean2)]

# Create reduced dataset with the top predictors and target variable (log_price)
df_reduced <- df_clean2 %>% select(all_of(top_20_vars), log_price)

# View column names in the reduced dataset
cat("Variables in df_reduced:\n")
print(colnames(df_reduced))

...

```

```
### **Split this new reduced dataset**
```

```
- **60/20/20 train-validation-test split**
```

```
```{r}
First extract the test dataset:
set.seed(123)
n <- nrow(df_reduced)
test_index <- sample(1:n, size = 0.2 * n)
test_reduced <- df_reduced[test_index,]
train_val <- df_reduced[-test_index,]

Then split the rest into validation and training sets:
val_index <- sample(1:nrow(train_val), size = 0.25 * nrow(train_val)) # 25% of 80% = 20%
valid_reduced <- train_val[val_index,]
train_reduced <- train_val[-val_index,]
```
```

```
### 2. Random Forest (on log_price)
```

Tune Hyperparameters (mtry and ntree) Using the Validation Set

```
```{r}
RMSE evaluation function
evaluate_model <- function(pred, actual) {
 mse <- mean((actual - pred)^2)
 rmse <- sqrt(mse)
 return(list(MSE = mse, RMSE = rmse))
}
```
```

Fit the Final Random Forest on the Full Training Set with Best Parameters

```
```{r}
Manual tuning of mtry and ntree
mtry_values <- c(10, 15, 20)
ntree_values <- c(100, 300, 500)
tune_results <- expand.grid(mtry = mtry_values, ntree = ntree_values)
tune_results$MSE <- NA
tune_results$RMSE <- NA

for (i in 1:nrow(tune_results)) {
 rf_temp <- randomForest(log_price ~ ., data = train_reduced,
 mtry = tune_results$mtry[i],
 ntree = tune_results$ntree[i])
 pred_val <- predict(rf_temp, newdata = valid_reduced)
 eval <- evaluate_model(exp(pred_val), exp(valid_reduced$log_price))
 tune_results$MSE[i] <- eval$MSE
}
```

```

 tune_results$RMSE[i] <- eval$RMSE
}

tune_results <- tune_results[order(tune_results$RMSE),]
print(tune_results)
best_combo <- tune_results[1,]
```

```{r}
Final RF model using best tuning parameters
rf_final <- randomForest(log_price ~ ., data = train_reduced,
 mtry = best_combo$mtry,
 ntree = best_combo$ntree)

Validation set predictions and evaluation
pred_val_final <- predict(rf_final, newdata = valid_reduced)
eval_val_final <- evaluate_model(exp(pred_val_final), exp(valid_reduced$log_price))
val_mse <- mean((exp(pred_val_final) - exp(valid_reduced$log_price))^2)

cat("Validation RMSE: $", round(eval_val_final$RMSE, 2), "\n")
cat("Validation MSE:", round(val_mse, 2), "\n\n")

Test set predictions and evaluation
pred_test_final <- predict(rf_final, newdata = test_reduced)
eval_test_final <- evaluate_model(exp(pred_test_final), exp(test_reduced$log_price))
test_mse <- mean((exp(pred_test_final) - exp(test_reduced$log_price))^2)

cat("Test RMSE: $", round(eval_test_final$RMSE, 2), "\n")
cat("Test MSE:", round(test_mse, 2), "\n")
```

**Lasso Regression (on log_price)**

Lasso Regression ( $\alpha = 1$ )

```{r}
Build model matrix once from full reduced dataset
x_full <- model.matrix(log_price ~ ., data = df_reduced)[, -1]
y_full <- df_reduced$log_price

Perform consistent 60/20/20 split on full data
set.seed(123)
n <- nrow(df_reduced)
test_index <- sample(1:n, size = 0.2 * n)
train_val_index <- setdiff(1:n, test_index)

val_index <- sample(train_val_index, size = 0.25 * length(train_val_index))

```

```

train_index <- setdiff(train_val_index, val_index)

Subset model matrix and target vectors
x_train <- x_full[train_index,]
x_valid <- x_full[val_index,]
x_test <- x_full[test_index,]

y_train <- y_full[train_index]
y_valid <- y_full[val_index]
y_test <- y_full[test_index]

```

```{r}
Cross-validation for best lambda
set.seed(123)
lasso_cv <- cv.glmnet(x_train, y_train, alpha = 1, standardize = TRUE)

Plot CV curve
plot(lasso_cv)

Fit final model
best_lambda_lasso <- lasso_cv$lambda.min
cat("Best lambda (Lasso):", best_lambda_lasso, "\n")

lasso_model <- glmnet(x_train, y_train, alpha = 1, lambda = best_lambda_lasso)

Predict
pred_valid_lasso <- predict(lasso_model, s = best_lambda_lasso, newx = x_valid)
pred_test_lasso <- predict(lasso_model, s = best_lambda_lasso, newx = x_test)

Evaluate
eval_val_lasso <- evaluate_model(exp(pred_valid_lasso), exp(y_valid))
eval_test_lasso <- evaluate_model(exp(pred_test_lasso), exp(y_test))

Display both MSE and RMSE
cat("Lasso Validation MSE:", round(eval_val_lasso$MSE, 2), "\n")
cat("Lasso Validation RMSE: $", round(eval_val_lasso$RMSE, 2), "\n\n")

cat("Lasso Test MSE:", round(eval_test_lasso$MSE, 2), "\n")
cat("Lasso Test RMSE: $", round(eval_test_lasso$RMSE, 2), "\n")

Selected predictors
lasso_coefs <- coef(lasso_model)
selected_predictors <- rownames(lasso_coefs)[which(lasso_coefs != 0)]
cat("\nSelected predictors by Lasso:\n")
print(selected_predictors)

```

```

...

GLM (Gaussian with identity link)

prepossessing\

```{r}
# Convert percentage-like variables to numeric
percentage_vars <- c("host_acceptance_rate", "host_response_rate")

convert_percentages_to_numeric <- function(df) {
  for (var in percentage_vars) {
    if (var %in% colnames(df)) {
      df[[var]] <- as.numeric(as.character(df[[var]]))
    }
  }
  return(df)
}

train_reduced <- convert_percentages_to_numeric(train_reduced)
valid_reduced <- convert_percentages_to_numeric(valid_reduced)
test_reduced <- convert_percentages_to_numeric(test_reduced)

# Remove factors with only one level
drop_single_level_factors <- function(df) {
  keep <- sapply(df, function(col) !(is.factor(col) && length(unique(col)) == 1))
  df[, keep]
}

train_reduced <- drop_single_level_factors(train_reduced)

# Align validation/test with training columns
valid_reduced <- valid_reduced[, colnames(train_reduced)]
test_reduced <- test_reduced[, colnames(train_reduced)]

# Drop NAs
train_reduced <- na.omit(train_reduced)
valid_reduced <- na.omit(valid_reduced)
test_reduced <- na.omit(test_reduced)

...

```{r}
Fit GLM on log_price
glm_log_model <- glm(log_price ~ ., data = train_reduced, family = gaussian(link = "identity"))

Predict on validation and test

```



```

pred_log_val <- predict(glm_log_model, newdata = valid_reduced)
pred_log_test <- predict(glm_log_model, newdata = test_reduced)

Actual and predicted prices (back-transform)
actual_val_price <- exp(valid_reduced$log_price)
actual_test_price <- exp(test_reduced$log_price)

pred_val_price <- exp(pred_log_val)
pred_test_price <- exp(pred_log_test)

MSE & RMSE
mse_price_val <- mean((actual_val_price - pred_val_price)^2)
mse_price_test <- mean((actual_test_price - pred_test_price)^2)

rmse_price_val <- sqrt(mse_price_val)
rmse_price_test <- sqrt(mse_price_test)

Output summary results
glm_metrics <- list(
 Validation_MSE = round(mse_price_val, 2),
 Validation_RMSE = round(rmse_price_val, 2),
 Test_MSE = round(mse_price_test, 2),
 Test_RMSE = round(rmse_price_test, 2)
)

print(glm_metrics)

...

```{r}
summary(glm_log_model)
```

XGBoost

Prepare Data (Convert to DMatrix)

```{r}
# Drop target from features
X_train <- model.matrix(log_price ~ . -1, data = train_reduced)
y_train <- train_reduced$log_price

X_valid <- model.matrix(log_price ~ . -1, data = valid_reduced)
y_valid <- valid_reduced$log_price

X_test <- model.matrix(log_price ~ . -1, data = test_reduced)

```

```

y_test <- test_reduced$log_price

# Convert to DMatrix
dtrain <- xgb.DMatrix(data = X_train, label = y_train)
dvalid <- xgb.DMatrix(data = X_valid, label = y_valid)
dtest <- xgb.DMatrix(data = X_test, label = y_test)

```

```

'''

```

Define Parameters and Train

```

'''{r}
# Hyperparameters
params <- list(
  objective = "reg:squarederror",
  eval_metric = "rmse",
  eta = 0.1,
  max_depth = 6
)

# Cross-validation to determine best nrounds
cv <- xgb.cv(
  params = params,
  data = dtrain,
  nrounds = 500,
  nfold = 5,
  early_stopping_rounds = 10,
  verbose = 0
)

best_nrounds <- cv$best_iteration

# Final model training using best_nrounds
xgb_model <- xgb.train(
  params = params,
  data = dtrain,
  nrounds = best_nrounds,
  watchlist = list(train = dtrain, eval = dvalid),
  early_stopping_rounds = 10,
  print_every_n = 20
)

```

```

# Save the model for later use (optional)
xgb.save(xgb_model, "final_xgb_model.model")
# Reload the model if needed
# xgb_model <- xgb.load("final_xgb_model.model")
'''

```

Predict on Validation & Test Sets

```
```{r}
Predict log prices
pred_valid_xgb <- predict(xgb_model, newdata = dvalid)
pred_test_xgb <- predict(xgb_model, newdata = dtest)

Convert back to price
pred_valid_price <- exp(pred_valid_xgb)
pred_test_price <- exp(pred_test_xgb)

actual_valid_price <- exp(y_valid)
actual_test_price <- exp(y_test)

Calculate RMSE and MSE
rmse_xgb_val <- sqrt(mean((actual_valid_price - pred_valid_price)^2))
rmse_xgb_test <- sqrt(mean((actual_test_price - pred_test_price)^2))

mse_xgb_val <- mean((actual_valid_price - pred_valid_price)^2)
mse_xgb_test <- mean((actual_test_price - pred_test_price)^2)

Output results
cat("XGBoost Validation RMSE: $", round(rmse_xgb_val, 2), "\n")
cat("XGBoost Validation MSE: ", round(mse_xgb_val, 4), "\n\n")

cat("XGBoost Test RMSE: $", round(rmse_xgb_test, 2), "\n")
cat("XGBoost Test MSE: ", round(mse_xgb_test, 4), "\n")

```
```

Feature Importance Plot

```
```{r}
Feature importance plot
importance_matrix <- xgb.importance(model = xgb_model)
xgb.plot.importance(importance_matrix, top_n = 20)
```
```

Gradient Boosted Trees (GBM)

```
```{r}
Prepare the data
gbm requires the response to be in the original (not log) scale
train_gbm <- train_reduced
valid_gbm <- valid_reduced
test_gbm <- test_reduced
```

```

train_gbm$price <- exp(train_gbm$log_price)
valid_gbm$price <- exp(valid_gbm$log_price)
test_gbm$price <- exp(test_gbm$log_price)

Remove the log_price column
train_gbm$log_price <- NULL
valid_gbm$log_price <- NULL
test_gbm$log_price <- NULL

Fit the GBM model

set.seed(123)
gbm_model <- gbm(
 formula = price ~ .,
 distribution = "gaussian",
 data = train_gbm,
 n.trees = 1000,
 interaction.depth = 4,
 shrinkage = 0.01,
 n.minobsinnode = 10,
 cv.folds = 5,
 verbose = FALSE
)

Determine the best number of trees
best_iter <- gbm.perf(gbm_model, method = "cv")

Predict
pred_valid_gbm <- predict(gbm_model, newdata = valid_gbm, n.trees = best_iter)
pred_test_gbm <- predict(gbm_model, newdata = test_gbm, n.trees = best_iter)

actual_valid_gbm <- valid_gbm$price
actual_test_gbm <- test_gbm$price

Evaluate
rmse_gbm_val <- sqrt(mean((actual_valid_gbm - pred_valid_gbm)^2))
rmse_gbm_test <- sqrt(mean((actual_test_gbm - pred_test_gbm)^2))

mse_gbm_val <- mean((actual_valid_gbm - pred_valid_gbm)^2)
mse_gbm_test <- mean((actual_test_gbm - pred_test_gbm)^2)

cat("GBM Validation RMSE: $", round(rmse_gbm_val, 2), "\n")
cat("GBM Validation MSE: ", round(mse_gbm_val, 4), "\n\n")

cat("GBM Test RMSE: $", round(rmse_gbm_test, 2), "\n")
cat("GBM Test MSE: ", round(mse_gbm_test, 4), "\n")

```

```
...
```

```
Q6
```

Diagnostics & Validation

```
```{r}
```

```
# Residuals on test set (actual - predicted)
```

```
residuals_test <- actual_test_price - pred_test_price
```

```
# Plot residuals vs predicted prices
```

```
library(ggplot2)
```

```
ggplot(data = data.frame(predicted = pred_test_price, residuals = residuals_test),
```

```
  aes(x = predicted, y = residuals)) +
```

```
geom_point(alpha = 0.5, color = "#1f77b4") +
```

```
geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
```

```
labs(
```

```
  title = "Residuals vs Predicted Price (Test Set)",
```

```
  x = "Predicted Price ($)",
```

```
  y = "Residuals ($)"
```

```
) +
```

```
theme_minimal()
```

```
...
```

```
```{r}
```

```
Standardized residuals on the test set
```

```
standardized_residuals <- residuals_test / sd(residuals_test)
```

```
Plot standardized residuals to spot potential outliers
```

```
library(ggplot2)
```

```
ggplot(data = data.frame(predicted = pred_test_price, std_resid = standardized_residuals),
```

```
 aes(x = predicted, y = std_resid)) +
```

```
geom_point(alpha = 0.5, color = "#2ca02c") +
```

```
geom_hline(yintercept = c(-2, 0, 2), linetype = c("dashed", "solid", "dashed"), color = c("red", "black", "red")) +
```

```
labs(
```

```
 title = "Standardized Residuals vs Predicted Price",
```

```
 x = "Predicted Price ($)",
```

```
 y = "Standardized Residuals"
```

```
) +
```

```
theme_minimal()
```

```
...
```

Q7 in the report

Question 8: Final prediction on test\_data using trained XGBoost model

```

```{r}
test_data <- read_excel("/Users/audreychang/Desktop/ACTL_ML/Textbook_R/HW/Assignment2/AirbnbTest (1).xlsx")
```

```

## Preprocess Test Data

```

```{r}
#days_between_reviews
test_data$days_between_reviews <- as.numeric(as.Date(test_data$last_review) - as.Date(test_data$first_review))
#drop first_review, last_review
test_data <- subset(test_data, select = -c(first_review, last_review))

```

#only run sucessfully when u have the original file

```

```

```

```

```{r}
#host_experience_years
test_data$host_experience_years <- 2025 - as.numeric(format(as.Date(test_data$host_since), "%Y"))
#drop host_since
test_data <- subset(test_data, select = -c(host_since))
```

```

```

```{r}
#avg_score
test_data$avg_score <- rowMeans(test_data[, c("review_scores_rating", "review_scores_accuracy",
"review_scores_cleanliness", "review_scores_checkin", "review_scores_communication", "review_scores_location",
"review_scores_value")], na.rm = TRUE)
#drop review_related variables
test_data <- test_data %>%
  select(-review_scores_rating,
        -review_scores_accuracy,
        -review_scores_cleanliness,
        -review_scores_checkin,
        -review_scores_communication,
        -review_scores_location,
        -review_scores_value)
```

```

```

```{r}
#Convert 't'/'f' to binary (0/1)
binary_vars <- c("host_is_superhost", "host_identity_verified", "host_has_profile_pic", "instant_bookable")
for (var in binary_vars) {
  test_data[[var]] <- ifelse(test_data[[var]] == "t", 1, 0)
}
```

```

```

```{r}
#numeric conversions
test_data <- test_data %>%
  mutate(
    host_listings_count = as.numeric(host_listings_count),
    host_total_listings_count = as.numeric(host_total_listings_count),
    host_response_rate = as.numeric(host_response_rate),
    host_acceptance_rate = as.numeric(host_acceptance_rate)
  )
```

```

## Handling NA

```

```{r}
#Identify columns in test_data where "N/A" appears (as a string)
na_string_cols_test <- sapply(test_data, function(x) {
  if (is.character(x) || is.factor(x)) {
    any(x == "N/A", na.rm = TRUE)
  } else {
    FALSE
  }
})

# List of column names that contain "N/A" strings
cols_to_fix_test <- names(na_string_cols_test[na_string_cols_test == TRUE])

#Replace "N/A" with real NA only in those columns
test_data[cols_to_fix_test] <- lapply(test_data[cols_to_fix_test], function(x) na_if(x, "N/A"))

#Count NA values per column
na_count_test <- sapply(test_data, function(x) sum(is.na(x)))
na_count_test <- na_count_test[na_count_test > 0] # Only show columns with missing
na_count_test
```

```

remove unused column

```

```{r}
test_data <- test_data %>%
  select(-bathrooms_text)
```

```{r}
# Convert percentage columns to numeric
test_data <- test_data %>%
  mutate(

```

```

    host_response_rate = as.numeric(host_response_rate),
    host_acceptance_rate = as.numeric(host_acceptance_rate)
  )

# Impute missing values in host_response_rate with its median
test_data$host_response_rate[is.na(test_data$host_response_rate)] <- median(test_data$host_response_rate, na.rm =
TRUE)

# Impute missing values in host_acceptance_rate with its median
test_data$host_acceptance_rate[is.na(test_data$host_acceptance_rate)] <- median(test_data$host_acceptance_rate, na.rm
= TRUE)

...

```{r}
Replace missing values for host_response_rate and host_acceptance_rate host_response_time, host_is_superhost with
"Unknown"

test_data$host_response_time[is.na(test_data$host_response_time)] <- "Unknown"
test_data$host_is_superhost[is.na(test_data$host_is_superhost)] <- "Unknown"
...

```{r}
# Check if there are any NA values in the entire test_data
anyNA(test_data)
...

```{r}
#Region Classification (based on neighbourhood)
test_data$region <- case_when(
 test_data$neighbourhood %in% east ~ "East",
 test_data$neighbourhood %in% west ~ "West",
 test_data$neighbourhood %in% north ~ "North",
 test_data$neighbourhood %in% south ~ "South",
 TRUE ~ "Other"
)

...

```{r}
#Property Group Categorization
test_data$property_group <- ifelse(test_data$property_type %in% c(
  "Entire rental unit", "Entire condo", "Entire serviced apartment",
  "Entire loft", "Room in serviced apartment", "Room in aparthotel",
  "Private room in serviced apartment", "Private room in condo"),
  "Apartment/Condo",

  ifelse(test_data$property_type %in% c(

```



```

# Apply the PCA transformation using the training PCA object
# Use predict() to project the test data onto the training PCA space
pca_capacity_test <- predict(pca_capacity, newdata = capacity_vars_test)

# Add the first principal component to the test data
test_data$PC_capacity <- pca_capacity_test[, 1]

# Drop the original capacity variables from the test set
test_data <- test_data %>% select(-accommodates, -beds, -bedrooms, -bathrooms)

'''

```

Apply the Model to the Test Set

```

'''{r}
colnames(test_data)
'''

```

3\. Make predictions using the trained model

```

'''{r}
# Assuming you have the list of top predictor variables from Question 5
test_data_reduced <- test_data %>% select(all_of(top_20_vars)) # Reduced dataset with top predictors

# Load the trained model from Question 5 (XGBoost in this case)
xgb_model <- xgb.load("final_xgb_model.model")

# Convert test data to DMatrix (for XGBoost)
dtest <- xgb.DMatrix(data = model.matrix(~ . -1, data = test_data_reduced)) # Ensure you remove the intercept term

# Predict log prices using the model
pred_log_price <- predict(xgb_model, newdata = dtest)

# Convert log price predictions back to the price scale (exponentiate)
test_data$price_prediction <- exp(pred_log_price) # Update the price_prediction column

# Optionally, check the first few predictions
head(test_data$price_prediction)

# Save the predictions to an Excel file
write_xlsx(test_data, "/Users/audreychang/Desktop/ACTL_ML/Textbook_R/HW/Assignment2/AirbnbTest (1).xlsx")

'''

```

Generative AI usage

OpenAI.(2025). *ChatGPT* (March 27 version) [Large language model]. <https://chat.openai.com/chat>

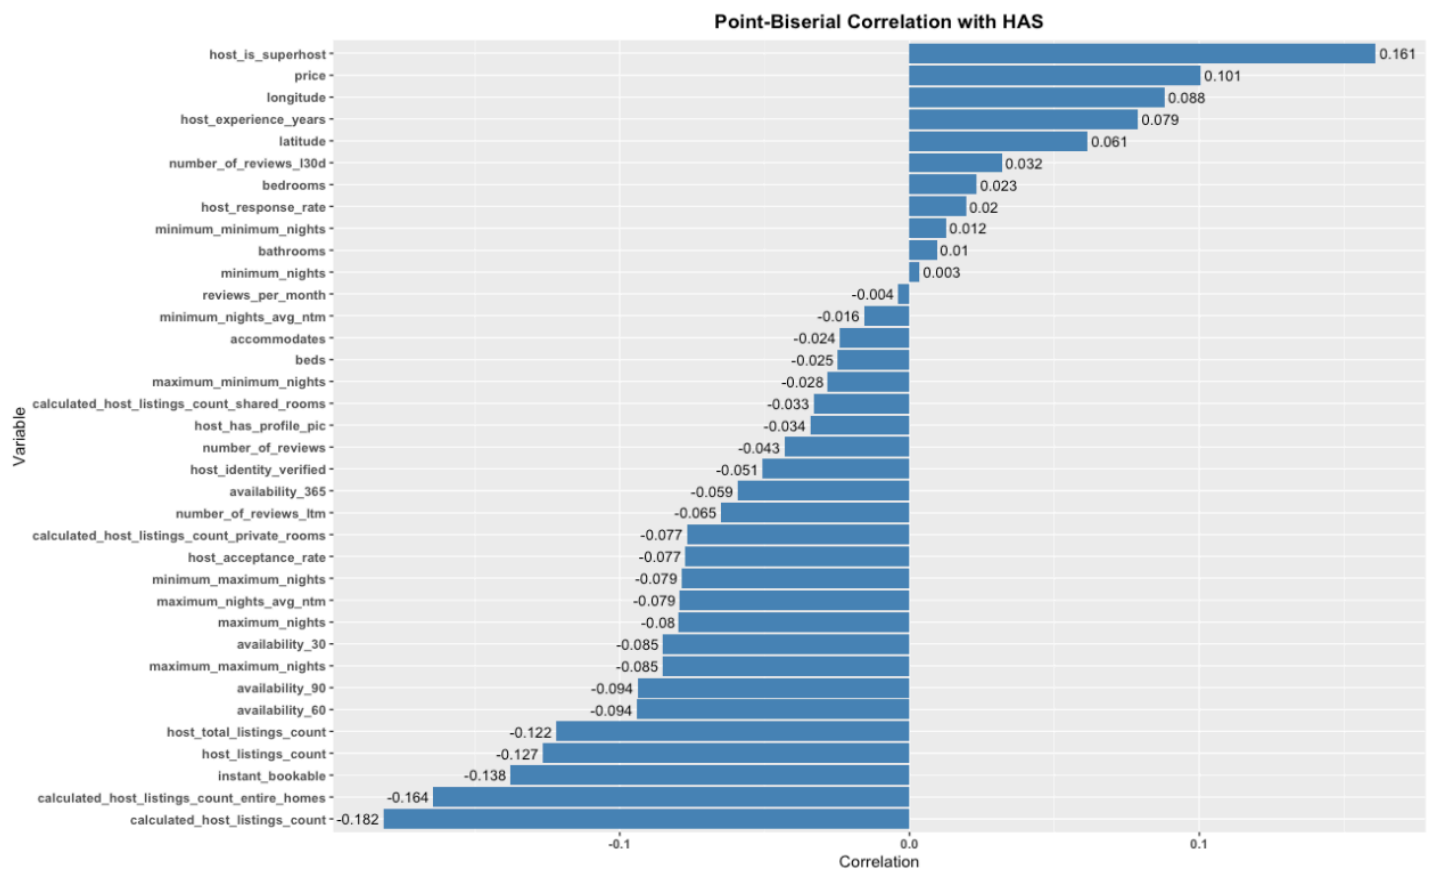
Purpose of Use:

ChatGPT was used to support various aspects of the assignment, including:

- **Editing and Refinement:**
 - Improved clarity, flow, and conciseness of report sections.
 - Reduced redundancy in model descriptions and interpretations.
- **Planning and Structure:**
 - Helped organize report sections logically.
 - Suggested concise formats for presenting model results and diagnostics.
- **Concept Clarification and Idea Generation:**
 - Explained statistical concepts such as:
 - Cook's Distance (influence diagnostics)
 - XGBoost mechanisms and hyperparameters
 - Gradient Boosted Trees and ensemble methods
 - Residual diagnostics in logistic regression and XGBoost
 - Suggested realistic improvement steps for the predictive model.
- **Prompt Examples Used:**
 - "Make this paragraph more concise but preserve meaning."
 - "Explain how to interpret Cook's Distance."
 - "What are the pros and cons of XGBoost in model interpretability?"
 - "Suggest ways to improve a regression model using advanced techniques."
 - "Refine this model limitation section to save words."

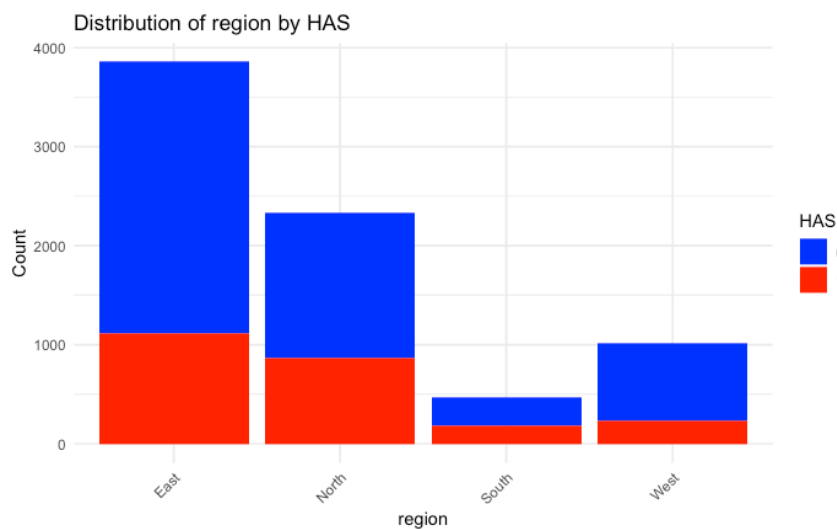
Others:

2.3 Variable Selection (Correlation plot)

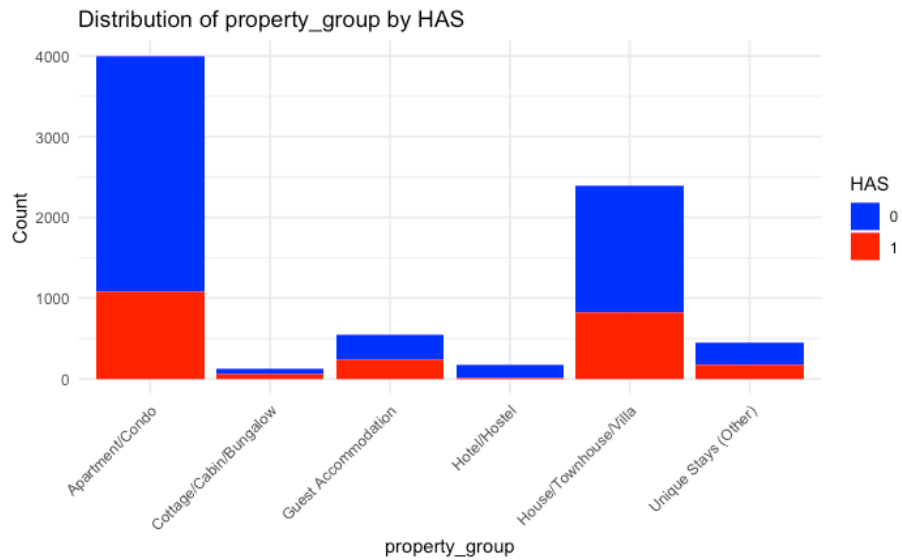


2.3 Data Visualization

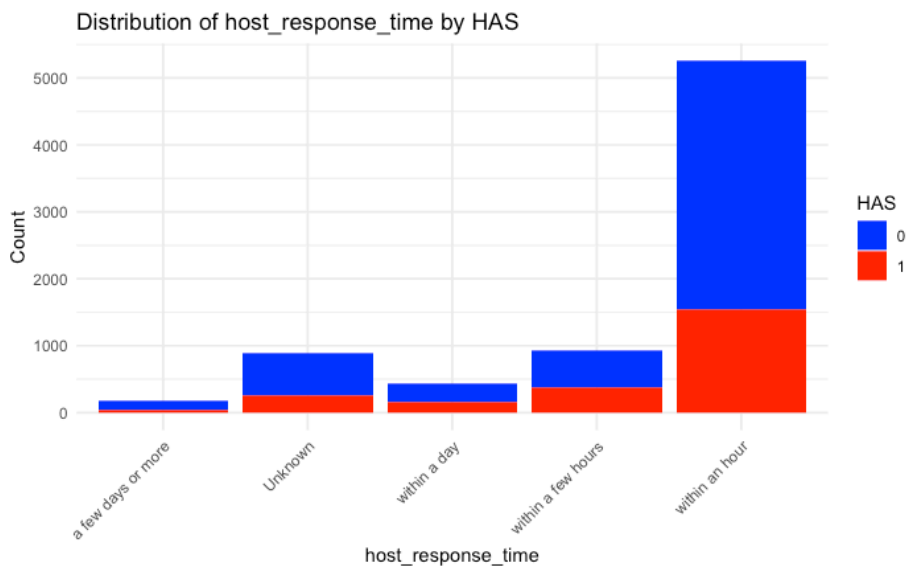
"Region"



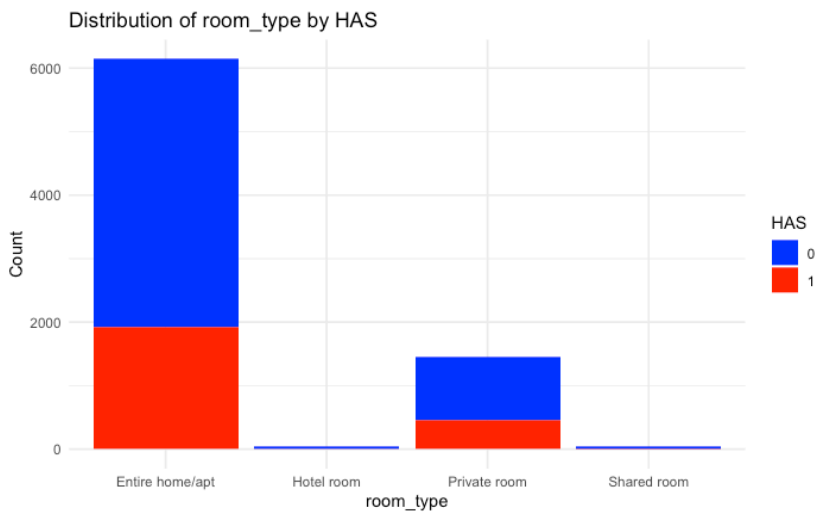
"property_group"



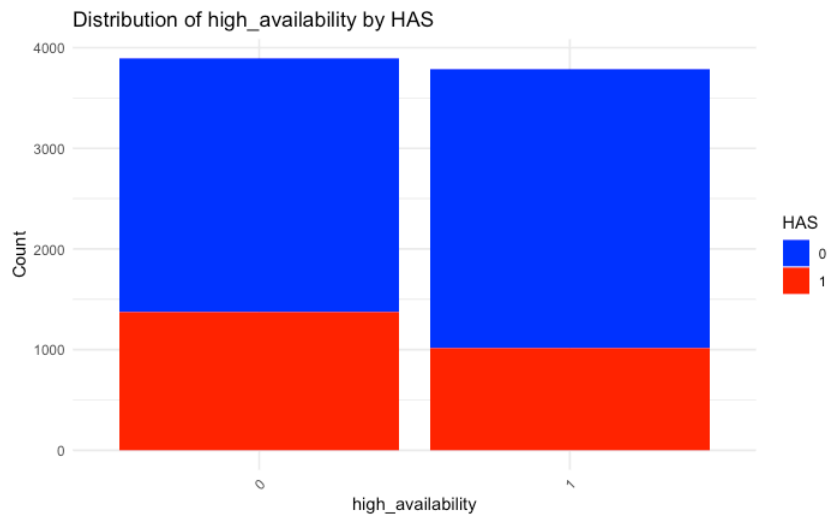
"Host_response_time"



"room_type"




High availability



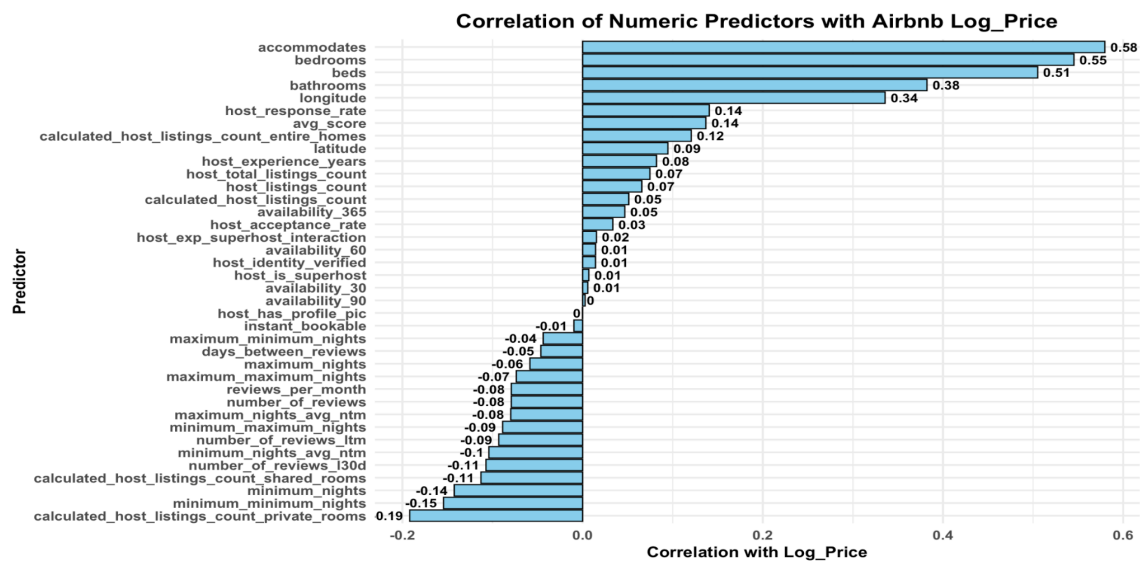
Predictors used in Q2

- host_is_superhost
- log(price
- longitude
- (host_experience_years
- latitude
- calculated_host_listings_count
- instant_bookable
- high_availability
- "Region"
- "property_group"
- "Host_response_time"
- "room_type"


Random Forest Variable Importance(Classification)



5.1 PCA




6.2 Raw Residual Analysis




5.2 Feature Selection Random Forest's %IncMSE metric

Variable Importance (%IncMSE)



5.3.2 Lasso Regression lambda.min (0.000197037)



5.3.4 XGBoost Feature Importance Plot

